

Linear regression using Stan

Alessandro Varacca, Thomas Heckelei

July 2024

But what does linear regression mean? (1)

$$y \sim \mathcal{D}(\theta)$$

- ▶ where θ indicates a vector of parameters;
- ▶ In standard **linear** regression models, we want to specify a functional form for $\mathbb{E}[y|X=x]$ **independently** of the (conditional) variance $\mathbb{V}(y|X=x)$;
- ▶ There are only a few choices of \mathcal{D} that allow to do so;
- ▶ The most popular one is the **Normal** distribution. Other options include the **Student's t** and the **Double exponential** distribution;
- ▶ The former is by far the most widely adopted;
- ▶ Therefore, $\theta = [\mu, \sigma]$ and:

$$y \sim \mathcal{N}(\mu, \sigma)$$

But what does linear regression mean? (2)

- For $i = 1 \dots, N$ independent observations, we typically write:

$$y_i \sim \mathcal{N}(\mu_i, \sigma)$$

$$\mu_i = \mathbb{E}[y|X = x] = \alpha + \sum_{p=1}^P \beta_p x_{p,i}$$

- or, equivalently:

$$y_i = \mu_i + \varepsilon_i$$

$$\mu_i = \mathbb{E}[y|X = x] = \alpha + \sum_{p=1}^P \beta_p x_{p,i}$$

$$\varepsilon_i \sim \mathcal{N}(0, \sigma^2)$$

- This is the classic homoscedastic linear regression model.

How many parameters?

- ▶ The classical linear regression model has $P + 2$ parameters:
 1. σ
 2. α
 3. all the β_p
- ▶ We therefore need to set a prior on each of these parameters;
- ▶ To understand what kind of priors we need, it is important to understand what each of these parameters governs in the distribution on y_i ;
- ▶ However, since y is binary, we are better off reasoning in terms of π ;
- ▶ In particular,
 1. $\sigma = \sqrt{\mathbb{V}(y)}$ governs the **spread** of y , marginally;
 2. $\alpha = \mathbb{E}[y|X = 0]$ governs the **location** of y , marginally;
 3. $\beta_p = \beta_p \approx \mathbb{E}[y|X_p = x_p + 1] - \mathbb{E}[y|X_p = x_p]$ describes the average change in y when we shift x_p by one unit.

Calibrating the prior(s) (1)

- ▶ If we have any knowledge about y_i , we can use this information to set reasonable prior distributions on each of these parameters;
- ▶ Going back to Thomas' example on yield, we might know what the expected yield of some crop might be, and how volatile this average value can be;
- ▶ In other words, we might have **prior** information on both α and σ ;
- ▶ Then, in case $P = 1$ and x_i represents, for example, soil moisture, then the literature (or common sense/agronomic expertise) can provide upper and lower bounds for the prior on β ;
- ▶ The same line of reasoning also apply in case of $P > 1$.

Calibrating the prior(s) (2)

- ▶ There are however cases in which such details are not or only partially known;
- ▶ Then how do we set reasonable priors?
- ▶ In the next slides we are going to explore a simulation-based calibration approach to solve this puzzle;
- ▶ The idea is to:
 1. choose a prior for σ , α and β_p and draw potential parameter values from these distribution;
 2. plug these parameter values into $N(\alpha + \sum_{p=1}^P \beta_p x_{p,i}, \sigma)$ and simulate potential values of y_i .

Load data and libraries

```
library(truncnorm)
library(rstan)


setwd("your working directory")

load("linear_regression_data.RData")
source("aux_functions.R")

set.seed(123)
```

Visualize the data (1)

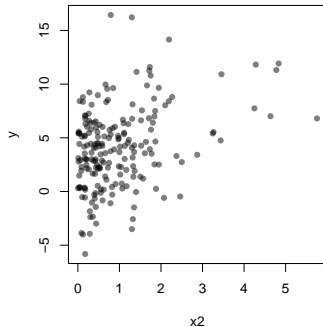
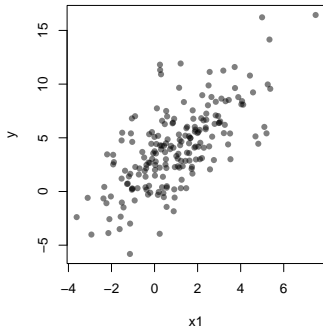
```
par(mfrow=c(1,3), pty="s")  
hist(y)  
hist(x1)  
hist(x2)
```



```
par(mfrow=c(1,1), pty="m")
```


Visualize the data (2)


```
par(mfrow=c(1,2), pty="s")  
plot(x1, y, pch = 16, col = rgb(0,0,0,.5))  
plot(x2, y, pch = 16, col = rgb(0,0,0,.5))
```



```
par(mfrow=c(1,1), pty="m")
```

Visualize the data (3)

```
par(pty="s")  
pairs(cbind(x1,x2,y), pch = 16, col = rgb(0,0,0,.5))
```



```
par(pty="m")
```

Begin sampling from the prior: σ (1)

- ▶ We begin by studying how σ may impact $\mathbf{y} = [y_1, \dots, y_N]$;
- ▶ Since the σ can only take up positive values, the **prior distribution** must satisfy this restriction;
- ▶ There are several options: **Exponential** distribution, **Gamma** distribution, **Truncated Normal** distribution and others;
- ▶ Here we use the Truncated Normal:

$$\sigma \sim N_+(0, \lambda)$$

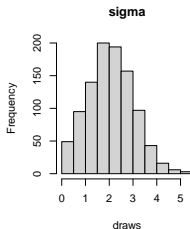
- ▶ Calibrating the prior for σ consist in choosing a sensible value for λ .

Begin sampling from the prior: σ (2)

- For example, for $\lambda = 2$, this is how 1000 draws from $N_+(0, 2)$ look like:

```
std_dev <- rtruncnorm(1000, 0, Inf, 2)
```

```
par(pty="s")  
hist(std_dev, main = "sigma", xlab = "draws")
```



```
par(pty="s")
```

Begin sampling from the prior: σ (3)

- ▶ We can now begin to study what $N_+(0, 2)$ means in terms of potential values of \mathbf{y} :
 1. Sample one parameter value from $N_+(0, \lambda)$, and get σ_y^s ;
 2. Sample one observation from $N(0, \sigma_y^s)$ and get $\tilde{\mathbf{y}}^s$;
 3. Repeat the process $S = 1000$ times.

```
y_sim <- matrix(nrow=1000, ncol=200)
for (s in 1:1000) {

  std_dev <- rtruncnorm(1, 0, Inf, 2)
  y_sim[s, ] <- rnorm(200, 0, std_dev)

}
```

- ▶ The $N \times S$ matrix $\mathbf{y_sim} = [\mathbf{y}^1, \dots, \mathbf{y}^S]$ approximates the so-called **Prior Predictive Distribution** (PrPD) of y ;
- ▶ We can then use visualization to study the properties of the PrPD and check whether our assumption makes any sense in face of the observed data, \mathbf{y} .

Visualizing the PrPD (1)

1. For each draw $\tilde{\mathbf{y}}^s$ calculate $\max^s = \max(\tilde{\mathbf{y}}^s)$;
2. For each draw $\tilde{\mathbf{y}}^s$ calculate $\min^s = \min(\tilde{\mathbf{y}}^s)$;
3. For each draw $\tilde{\mathbf{y}}^s$ calculate $\text{sd}^s = \text{sd}(\tilde{\mathbf{y}}^s)$

```
max_y <- apply(y_sim, 1, max)
min_y <- apply(y_sim, 1, min)
sd_y <- apply(y_sim, 1, sd)
```

These are obviously $S \times 1$ vectors:

```
length(max_y)
```

```
## [1] 1000
```

```
length(min_y)
```

```
## [1] 1000
```


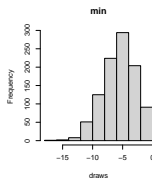
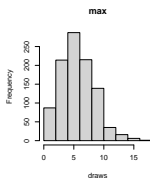
```
length(sd_y)
```

```
## [1] 1000
```

Visualizing the PrPD (2)

Plot these quantities:

```
par(mfrow=c(1,3), pty="s")  
hist(max_y, main = "max", xlab = "draws")  
hist(min_y, main = "min", xlab = "draws")  
hist(sd_y, main = "sd", xlab = "draws")
```






```
par(mfrow=c(1,1), pty="m")
```

Each histogram represents the extent of maximum/minimum values, and standard deviations that can be obtained setting $\lambda = 2$.

Visualizing the PrPD (3)

We can check consistency between the summary statistics of the PrPD and our observed data points by overlaying $\min(\mathbf{y})$, $\max(\mathbf{y})$ and $sd(\mathbf{y})$ to the corresponding histogram:

```
par(mfrow=c(1,3), pty="s")
hist(max_y, main = "max", xlab = "draws")
abline(v=max(y), col="red", lwd=2)
hist(min_y, main = "min", xlab = "draws")
abline(v=min(y), col="blue", lwd=2)
hist(sd_y, main = "sd", xlab = "draws")
abline(v=sd(y), col="green", lwd=2)
```



```
par(mfrow=c(1,1), pty="m")
```


Visualizing the PrPD (3)

- ▶ The standard deviation (SD) of the PrPD is smaller than the observed SD;
- ▶ Also the maximum value does not quite match;
- ▶ We can try changing *lambda* to, say, 3:




```
y_sim <- matrix(nrow=1000, ncol=200)
for (s in 1:1000) {

  std_dev <- rtruncnorm(1, 0, Inf, 3)
  y_sim[s, ] <- rnorm(200, 0, std_dev)

}
```

Visualizing the PrPD (4)

```
par(mfrow=c(1,3), pty="s")
hist(apply(y_sim, 1, max), main = "max", xlab = "draws")
abline(v=max(y),col="red", lwd=2)
hist(apply(y_sim, 1, min), main = "min", xlab = "draws")
abline(v=min(y),col="blue", lwd=2)
hist(apply(y_sim, 1, sd), main = "sd", xlab = "draws")
abline(v=sd(y),col="green", lwd=2)
```



```
par(mfrow=c(1,1), pty="m")
```

- ▶ SD is now well calibrated, but to have better max and min we will need to tweak α and β .

Calibrating α (1)

- ▶ Recall that $\alpha = \mathbb{E}[y|X = 0]$;
- ▶ Therefore, the range of plausible value of α should reflect the expected range of average values for y ;
- ▶ Since y includes both positive and negative values, using a **Normal** prior for α should do;
- ▶ We begin with $\alpha \sim N(0, 5)$;
- ▶ Notice that this prior implies that the range of $\mathbb{E}[y|X = 0]$ is roughly $[-15, +15]$ and values between $[-5, 5]$ are more likely.

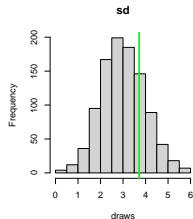
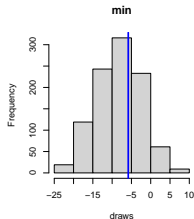
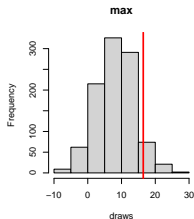
```
y_sim <- matrix(nrow=1000, ncol=200)
for (s in 1:1000) {

  std_dev <- rtruncnorm(1, 0, Inf, 3)
  mu <- rnorm(1, 0, 5)
  y_sim[s, ] <- rnorm(200, mu, std_dev)

}
```

Calibrating α (2)

```
par(mfrow=c(1,3), pty="s")
hist(apply(y_sim, 1, max), main = "max", xlab = "draws")
abline(v=max(y),col="red", lwd=2)
hist(apply(y_sim, 1, min), main = "min", xlab = "draws")
abline(v=min(y),col="blue", lwd=2)
hist(apply(y_sim, 1, sd), main = "sd", xlab = "draws")
abline(v=sd(y),col="green", lwd=2)
```

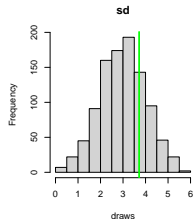
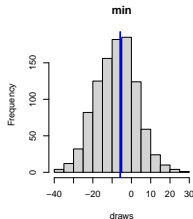



```
par(mfrow=c(1,1), pty="m")
```

Calibrating α (3)

- ▶ Setting α to $N(0, 10)$, instead leads to:

```
par(mfrow=c(1,3), pty="s")
hist(apply(y_sim, 1, max), main = "max", xlab = "draws")
abline(v=max(y), col="red", lwd=2)
hist(apply(y_sim, 1, min), main = "min", xlab = "draws")
abline(v=min(y), col="blue", lwd=2)
hist(apply(y_sim, 1, sd), main = "sd", xlab = "draws")
abline(v=sd(y), col="green", lwd=2)
```



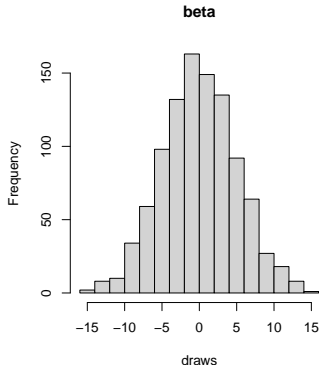
```
par(mfrow=c(1,1), pty="m")
```

Calibrating α (3)

- ▶ Observe that, although these priors might seem at first quite informative, they are not!
- ▶ In fact, although the most likely values under the priors set so far are close to those observed in the dataset, the range of plausible y extends quite beyond that range;
- ▶ Prior distributions of this nature are typically referred to as **weakly informative**;
- ▶ There is a very large chunk of the literature suggesting that weakly informative priors should be preferred to completely uninformative ones;
- ▶ There are several reasons for this, but one very important one in probabilistic programming is that they usually trigger computational issues;
- ▶ Examples of uninformative priors that should not be used include the historically very popular $\text{Uniform}(-\infty, +\infty)$.

Calibrating β_1 and β_2 (1)

- Recall that $\beta_p \approx \mathbb{E}[y|X_p = x_p + 1] - \mathbb{E}[y|X_p = x_p]$. Since $\bar{y} = 4.07$, setting $\beta_p \sim N(0, 5)$ implies that shifting x_p by one unit would change the expected value of y by:



Calibrating β_1 and β_2 (2)

- ▶ Notably, by using a normal prior centered at zero, we are implicitly assuming that the most likely change in $\mathbb{E}[y|X_p = x_p]$ **before seeing the data** is small;
- ▶ This is a **conservative** choice and is typically adopted in regression problems;
- ▶ However, notice also that the $N(0, 5)$ prior allows the ‘slopes’ to potentially extend up to roughly $[-15, +15]$ and beyond;
- ▶ It is the likelihood-prior interaction that will cause a change in the **posterior** distribution of β_p .

Calibrating β_1 and β_2 (3)

- ▶ To simulate the impact of

1. $\sigma \sim N_+(0, 3)$
2. $\alpha \sim N(0, 10)$
3. $\beta_1 = \beta_2 \sim N(0, 5)$

on the PrPD of y , I have prepared a small Stan script;

- ▶ When we need to calibrate more than one prior, wiring down the whole data generating process in Stan is a convenient way to avoid hardcoding the entire scheme in R;
- ▶ In the next slides I will go through this program very briefly, just so you have an understanding of how it works.

Stan program for prior calibration (1)

- ▶ We first load the script using the R function `stan_model()`:

```
rstan_options(auto_write = TRUE)
lin_reg_prior <- stan_model("linear_regression_prior.stan")
```

- ▶ Before using using the `lin_reg_prior` object to simulate the data, let us have a look at what this code looks like!

Stan program for prior calibration (2)

- ▶ The input to **any** Stan program is a data **list**;
- ▶ This list contains all the (hyper)parameter values, the data and any additional variable necessary to run the code contained in the script

```
dat_list <- list(  
  N = length(y),  
  x1 = x1,  
  x2 = x2,  
  y = y,  
  mu_b1 = 0,  
  mu_b2 = 0,  
  sigma_b1 = 5,  
  sigma_b2 = 5,  
  mu_alpha = 0,  
  sigma_alpha = 10,  
  lambda_sigma = 3  
)
```

Stan program for prior calibration (3)

- ▶ N indicates the sample size;
- ▶ x_1 and x_2 are $N \times 1$ vector for the two covariates;
- ▶ y indicates the $N \times 1$ vector \mathbf{y} ;
- ▶ μ_{b1} and μ_{b2} indicate the mean of the priors on β_1 and β_2 , respectively;
- ▶ σ_{b1} and σ_{b2} indicate the standard deviation of the priors on β_1 and β_2 , respectively;
- ▶ μ_{α} indicates the mean of the prior on α ;
- ▶ σ_{α} indicates the standard deviation of the prior on α ;
- ▶ λ_{σ} indicates the standard deviation of the prior on σ .

Stan program for prior calibration (4)

- ▶ We can now simulate $\tilde{\mathbf{y}}^s$!

```
prior_sample <- sampling(lin_reg_prior,  
                        data = dat_list,  
                        algorithm = "Fixed_param")
```

- ▶ and extract the $N \times S$ matrix of simulated outcome values:



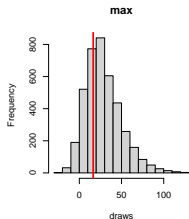
```
prior_sample_fit <- extract(prior_sample)  
y_sim <- prior_sample_fit$y_pred
```

- ▶ We can next repeat the visual assessment introduced in the earlier slides. We begin by calculating our summary statistics:

```
max_y <- apply(y_sim, 1, max)  
min_y <- apply(y_sim, 1, min)  
sd_y <- apply(y_sim, 1, sd)
```

Calibrated PrPD (1)

```
par(mfrow=c(1,3), pty="s")  
hist(max_y, main = "max", xlab = "draws")  
abline(v=max(y),col="red", lwd=2)  
hist(min_y, main = "min", xlab = "draws")  
abline(v=min(y),col="blue", lwd=2)  
hist(sd_y, main = "sd", xlab = "draws")  
abline(v=sd(y),col="green", lwd=2)
```





```
par(mfrow=c(1,1), pty="m")
```

Calibrated PrPD (2)

- ▶ We can also visualize the bundle of regression lines implied by $\beta_1 = \beta_2 \sim N(0, 5)$;
- ▶ To do this we can use the

```
par(mfrow=c(1,2))
plot_prior_lines(x1,
                 prior_sample_fit$alpha,
                 prior_sample_fit$beta1,
                 prior_sample_fit$y_pred,
                 ylim = c(-40, 40), xlab = "x1")
plot_prior_lines(x2,
                 prior_sample_fit$alpha,
                 prior_sample_fit$beta2,
                 prior_sample_fit$y_pred,
                 ylim = c(-40, 40), xlab = "x2")
par(mfrow=c(1,1))
```

Calibrated PrPD (3)



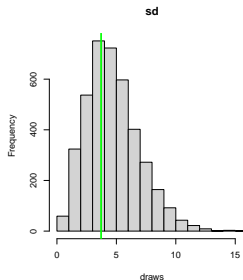
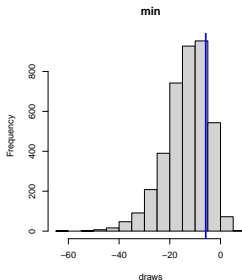
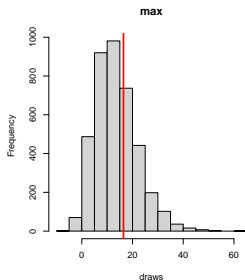
Calibrated PrPD (4)

- ▶ Things to observe:
 1. This model is general. Perhaps too general. . . It implies very extreme modelling assumption (i.e., potentially nonsensical regression lines!);
 2. Conditional on the two covariates x_1 and x_2 , the prior on the intercept maybe also too wide (look at how wide the bundle of regression lies is around zero!);
- ▶ Let us see what changes by adjusting both distributions!



Calibrated PrPD (5)

- ▶ Adjusting β_1 and β_2 to $N(0, 2)$, α to $N(0, 2.5)$ and re-running the simulation yields:

```
dat_list$sigma_b1 <- 2  
dat_list$sigma_b2 <- 2  
dat_list$sigma_alpha <- 2.5
```



Calibrated PrPD (6)



Default priors (1)

- ▶ This calibration exercise demonstrates that, when knowledge on y and x_p is limited, prior calibration can be tedious;
- ▶ Moreover, when the number of regressors increases and/or the model includes non-linearities and interactions, choosing the most suitable distribution for β_p becomes very hard;
- ▶ Also, new sample information could yield rather different (and bad!) PrPD if the distribution of the new y , x_1 and x_2 changes too much;
- ▶ A commonly adopted solution to this issue consists of **standardizing** our variables;
- ▶ These transformation remove scale-dependence and allow to defined one-fits-all priors on the parameters of interest;
- ▶ Let's see how!

Default priors (2)

- ▶ Standardizing some numeric variable z means:

$$\dot{z} = \frac{z - \bar{z}}{\text{sd}(z)}$$

- ▶ Therefore, any \dot{z} will have mean zero and standard deviation 1 by construction;
- ▶ Applying this straightforward transformation to our data simplifies enormously the calibration procedure discussed above;
- ▶ To begin with, by standardizing y , we force its mean to zero, so $\alpha = 0$ by construction;
- ▶ We therefore do not need to calibrate the prior for α anymore;
- ▶ Moreover, since $\mathbb{V}(\dot{y}) = 1$, setting $\sigma \sim N_+(0, 1)$ already provides a good default for the prior on the standard deviation!

Default priors (3)

- ▶ Finally, since $\mathbb{V}(\dot{y}) = 1$ and the range of \dot{y} is artificially constraint to roughly $[-3, +3]$ (three SD of \dot{y}), giving both β_1 and β_2 a $N(0, 1)$ prior seems a sensible choice;
- ▶ To see what these new values yield in terms of PrPD, let us redefine our data list accordingly:

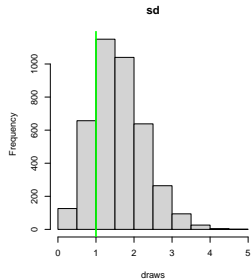
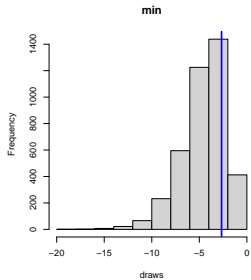
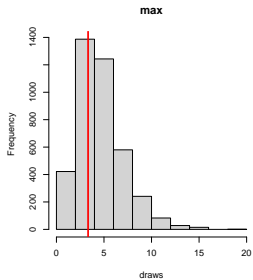
```
y_std <- as.vector(scale(y))
x1_std <- as.vector(scale(x1))
x2_std <- as.vector(scale(x2))

dat_list <- list(
  N = length(y_std),
  x1 = x1_std, x2 = x2_std,
  y = y_std,
  mu_b1 = 0,
  mu_b2 = 0,
  sigma_b1 = 1, sigma_b2 = 1,
  mu_alpha = 0, sigma_alpha = 0,
  lambda_sigma = 1
)
```



Default priors (4)

```
prior_sample <- sampling(lin_reg_prior,  
                        data = dat_list,  
                        algorithm = "Fixed_param")  
prior_sample_fit <- extract(prior_sample)  
y_sim <- prior_sample_fit$y_pred  
  
max_y <- apply(y_sim, 1, max)  
min_y <- apply(y_sim, 1, min)  
sd_y <- apply(y_sim, 1, sd)
```

Default priors (5)



Default priors (6)



Estimation

- ▶ Once all the priors have been setup, we can use Stan to sample from the implied joint posterior distribution:

$$f(\sigma, \beta_1, \beta_2 | \mathbf{y}, X_1 = x_1, X_2 = x_2)$$

- ▶ To do this, we need to load the second program:

```
lin_reg <- stan_model("linear_regression.stan")
```

- ▶ Fitting the model to the data now only requires a slightly different sampling statement:

```
fit <- sampling(lin_reg,  
               data = dat_list,  
               chains = 4,  
               cores = 4)  
sample_fit <- extract(fit)
```

Exploring the results (1)

To quickly visualize the results, we can call the function `summary` as we would do with the standard `lm` output:

```
params <- c("beta1", "beta2", "sigma")  
summary(fit, pars = params)$summary[,1:3]
```

```
##              mean      se_mean      sd  
## beta1 0.7648427 0.0007385151 0.04126889  
## beta2 0.4857550 0.0007458866 0.04067350  
## sigma 0.5618872 0.0004903962 0.02896625
```

```
params <- c("beta1", "beta2", "sigma")  
summary(fit, pars = params)$summary[,4:8]
```

```
##           2.5%      25%      50%      75%      97.5%  
## beta1 0.6822056 0.7378445 0.7646251 0.7917823 0.8481729  
## beta2 0.4045790 0.4586613 0.4854272 0.5127030 0.5666763  
## sigma 0.5080042 0.5416115 0.5608089 0.5806676 0.6213943
```

Exploring the results (2)

- ▶ Because of standardization, we need to transform all the coefficients back to the original scale of the data;
- ▶ Otherwise, $\hat{\beta}_p \approx \mathbb{E}[y|X_p = x_p + \text{sd}(y)] - \mathbb{E}[y|X_p = x_p]$;
- ▶ To do this, it suffices to rescale the standardized coefficients $\hat{\beta}_p$ by the standard deviation of x_p and y :

$$\beta_p = \hat{\beta}_p \times \frac{\text{sd}(y)}{\text{sd}(x_p)}$$


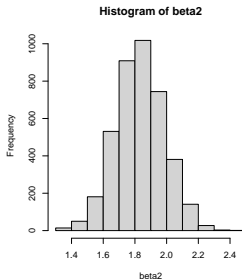
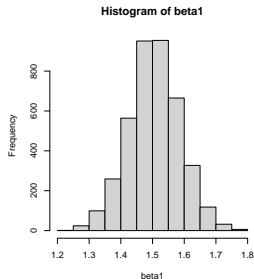
- ▶ Moreover, we can retrieve the intercept through:

$$\alpha = \bar{y} - \sum_{p=1}^P \beta_p \bar{x}_p$$

Exploring the results (3)

```
beta1 <- sample_fit$beta1 * (sd(y)/sd(x1))  
beta2 <- sample_fit$beta2 * (sd(y)/sd(x2))  
alpha <- mean(y)-(beta1*mean(x1))-(beta2*mean(x2))
```

```
par(mfrow=c(1,3), pty="s")  
hist(beta1)  
hist(beta2)  
hist(alpha)
```



```
par(mfrow=c(1,1), pty="m")
```

Exploring the results (4)

```
summary(beta1)
```

##	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
##	1.223	1.453	1.506	1.506	1.559	1.789



```
summary(beta2)
```

##	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
##	1.303	1.729	1.830	1.831	1.933	2.423

```
summary(alpha)
```

##	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
##	0.1201	0.6977	0.8171	0.8154	0.9368	1.4617

Exploring the results (5)



What if we used `lm()` instead?

```
summary(lm(y ~ x1 + x2))
```

```
##  
## Call:  
## lm(formula = y ~ x1 + x2)  
##  
## Residuals:  
##      Min       1Q   Median       3Q      Max   
## -5.6352 -1.4044  0.1549  1.3290  5.4982   
##  
## Coefficients:  
##              Estimate Std. Error t value Pr(>|t|)      
## (Intercept)  0.81122    0.23174    3.50 0.000574 ***  
## x1           1.50803    0.07953   18.96 < 2e-16 ***  
## x2           1.83350    0.15226   12.04 < 2e-16 ***  
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
##  
## Residual standard error: 2.082 on 197 degrees of freedom  
## Multiple R-squared:  0.689, Adjusted R-squared:  0.6859   
## F-statistic: 218.3 on 2 and 197 DF, p-value: < 2.2e-16
```


Leveraging the posterior

- ▶ So far, it might seem that all the extra effort to setup a probabilistic program to estimate a linear regression might be worthless in fact of a simple `lm()` call;
- ▶ However, having samples from the joint posterior unlocks several features that only belong in the Bayesian realm;
- ▶ First of all, it is worth stressing that having $f(\beta_p | \mathbf{y}, X = \mathbf{x})$ is different from having just $\mathbb{E}[\beta_p]$ and its standard error;
- ▶ If the parameters in $f(\sigma, \beta_1, \beta_2 | \mathbf{y}, X_1 = x_1, X_2 = x_2)$ are uncorrelated, then extracting any component from the $S \times (P + 2)$ matrix $[\boldsymbol{\sigma}, \boldsymbol{\alpha}, \beta_1, \beta_2]$, provides S samples from the marginal posterior of that parameter;
- ▶ These are the quantities that we are going to exploit.

Posterior probabilities (1)

- ▶ Consider the marginal posterior of β_2 , which we store in the vector `beta2`;
- ▶ We can use this object to calculate probabilities. For example, what is the probability that $\beta_2 > 2$?
- ▶ This can be easily calculated as the average count of posterior parameter values that are greater than 2:

```
thr <- 2  
prob <- mean(beta2>thr); prob
```

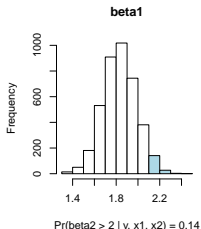
```
## [1] 0.13825
```

Posterior probabilities (2)

- We can also visualize this probability using the histogram above:

```
h_data <- hist(beta2, plot = F)
h_area <- cut(h_data$breaks, c(-Inf, thr, Inf))

par(pty="s")
plot(h_data, col = c("white", "lightblue")[h_area], main = "beta1",
     xlab = paste0("Pr(beta2 > ", thr, " | y, x1, x2) = ", round(prob,
```



```
par(pty="m")
```

Credible Intervals

- ▶ We can summarize the marginal posterior of our parameters via intervals;
- ▶ These are called Credible Intervals (Crl);
- ▶ They are naturally interpreted as the as the proportion of parameter values that fall between (i.e., what is the range of most likely values) an upper and lower value;
- ▶ Therefore, Crl are simply computed as the γ and $1 - \gamma$ **quantiles** of β_2 . For $\gamma = 0.025$ we have the 95% Crl:

```
quantile(beta2, c(.025, .975))
```

```
##      2.5%      97.5%
```

```
## 1.525164 2.136232
```

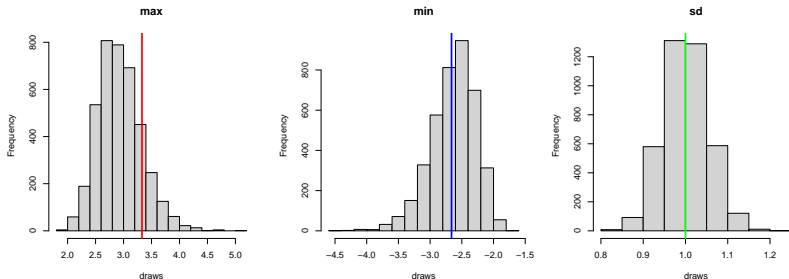
Model checking (internal consistency - 1)

- ▶ Another useful tool in the Bayesian regression toolbox is the sp-called Posterior Predictive Distribution (PoPD);
- ▶ Much like we used the PrPD to simulate y under different prior configurations, we use the PoPD to simulate y under the estimated posterior;
- ▶ The idea is to check the consistency between the simulated y and the observed counterpart;
- ▶ This automatically computed in the Stan program above and can be easily accessed by extracting `y_pred` from the fitted model:

```
y_pred <- sample_fit$y_pred
```

Model checking (internal consistency - 2)

```
par(mfrow=c(1,3), pty="s")  
hist(apply(y_pred, 1, max), main = "max", xlab = "draws")  
abline(v=max(y_std),col="red", lwd=2)  
hist(apply(y_pred, 1, min), main = "min", xlab = "draws")  
abline(v=min(y_std),col="blue", lwd=2)  
hist(apply(y_pred, 1, sd), main = "sd", xlab = "draws")  
abline(v=sd(y_std),col="green", lwd=2)
```



```
par(mfrow=c(1,1), pty="m")
```

Model checking (internal consistency - 3)

```
plot_prior_dens(y_pred, y_std, main = "PoPD")
```

