

Logistic regression using Stan

Alessandro Varacca, Thomas Heckelei

July 2024

But what does logistic regression mean? (1)

$$y \sim \mathcal{D}(\boldsymbol{\theta})$$

- ▶ where $\boldsymbol{\theta}$ indicates a vector of parameters;
- ▶ When the outcome of some variable y is binary, $\mathcal{D} = \text{Ber}(\pi)$, where:

$$\text{Ber}(\pi) = \begin{cases} \pi & \text{if } y = 1, \\ 1 - \pi & \text{if } y = 0. \end{cases}$$

and

$$\mathbb{E}[y] = \pi,$$

$$\mathbb{V}(y) = \pi(1 - \pi)$$

- ▶ In **logistic** regression models, we want to specify a functional form for $\mathbb{E}[y|X = x]$. However, unlike linear regression, this cannot be modelled **independently** of the (conditional) variance $\mathbb{V}(y|X = x)$;

But what does logistic regression mean? (2)

- ▶ Since $\pi \in [0, 1]$, when we specify a functional form $\mathbb{E}[y|X = x]$ we need to make sure that is is constrained between 0 and 1;
- ▶ This is typically done via the so-called **link functions**:

$$\mathbb{E}[y|X = x] = g(\mu) = g \left(\alpha + \sum_{p=1}^P \beta_p x_p \right)$$

where either

$$g(\mu) = \begin{cases} \Phi(\mu) & \text{for Probit regression} \\ \Lambda(\mu) = \frac{1}{1 + \exp(\mu)} & \text{for Logistic regression.} \end{cases}$$

- ▶ We will use the second one.

How many parameters?

- ▶ The logistic regression model has $P + 1$ parameters:
 1. α
 2. all the β_p
- ▶ We therefore need to set a prior on each of these parameters;
- ▶ To understand what kind of priors we need, it is important to understand what each of these parameters governs in the distribution on y_i ;
- ▶ However, since y is binary, it is typically preferable to reason in terms of π , instead;
- ▶ Moreover, any change in α or β_p does not impact π_i directly, but it does so through the transformation $g(\alpha)$ or $g(\beta_p)$;
- ▶ In other words,
$$\beta_p = g^{-1}\{\mathbb{E}[y|X_p = x_p + 1]\} - g^{-1}\{\mathbb{E}[y|X_p = x_p]\}$$
 and, because of that, they are hard to interpret on a probability scale;
- ▶ This makes prior calibration harder in logistic regression.

Calibrating the prior(s): α (1)

- ▶ We start off by studying a reasonable prior for α which, as in the linear regression case, is much easier to understand;
- ▶ Specifically, $\alpha = g^{-1}\{\mathbb{E}[y|X = 0]\} = g^{-1}\{\pi\}$;
- ▶ Therefore, to understand how a prior on α impacts π , we run the following simulation:
 1. Sample α^s from its prior;
 2. Calculate $\pi^s = g(\alpha^s)$
- ▶ Since this is relatively simple to implement, we immediately try different configurations:

$$\alpha \sim N(0, 10)$$

$$\alpha \sim N(0, 2.5)$$

$$\alpha \sim N(0, 1.5)$$

$$\alpha \sim N(0, 1)$$

Load data and libraries

```
library(rstan)
library(tidybayes)
library(tidyverse)
library(boot)


setwd("your working directory")

load("logistic_regression_data.RData")
source("aux_functions.R")

set.seed(123)
```

Visualize the data

```
par(mfrow=c(1,3), pty="s")  
hist(y)  
hist(x1)  
hist(x2)
```







```
par(mfrow=c(1,1), pty="m")
```

Calibrating the prior(s): α (2)

```
par(mfrow=c(1,4), pty="s")
for (j in c(1, 1.5, 2.5, 10)) {

  alpha <- rnorm(1000, 0, j)
  g_alpha <- inv.logit(alpha)
  hist(g_alpha, main = paste0("a ~ N(0, ", j, ")"), xlab = "pi")
  abline(v = 0.5, col = "red", lty = 2)

}
```



```
par(mfrow=c(1,1), pty="m")
```


Calibrating the prior(s): α (3)

- ▶ How to judge these priors?
- ▶ If we have no strong beliefs that π should be right, left skewed or skewed towards extreme ≈ 1 or ≈ 0 values, then $N(0, 2.5)$ and $N(0, 10)$ are **very** informative;
- ▶ In other words, choosing either of these distribution will yield priors placing a lot of emphasis on 'certain' zero or one values for y ;
- ▶ On the other hand, **no information** in the context of binary outcomes corresponds to $\pi \approx 0.5$;
- ▶ Put it differently, before seeing the data, any observation should have 50-50 chance of being one (or zero);
- ▶ Therefore, $\alpha \sim N(0, 1)$ is the most uninformative prior.

Calibrating β_1 and β_2 (1)

- ▶ Understanding what prior to give β_1 and β_2 is more complex;
- ▶ Not only do the distribution on these parameters will impact π , but so will their interaction with the covariates;
- ▶ As for α , let us see how different distribution choices for β_p play out: we begin by setting $\beta_1 = \beta_2 \sim N(0, 5)$ and simulate:
 1. α^s from $N(0, 1)$;
 2. β_1^s from $N(0, 5)$;
 3. β_2^s from $N(0, 5)$;
 4. $\pi_i^s = 1/(1 + \exp\{\alpha^s + \beta_1^s x_{i,1} + \beta_2^s x_{i,2}\})$.
- ▶ As in the linear regression case study, we run this simulation using a separate Stan program:

```
rstan_options(auto_write = TRUE)
logi_reg_prior <- stan_model("logistic_regression_prior.stan")
```

Calibrating β_1 and β_2 (2)

- As before, we must first set the data list:

```
dat_list <- list(  
  N = length(y),  
  x1 = x1,  
  x2 = x2,  
  y = y,  
  mu_b1 = 0,  
  mu_b2 = 0,  
  sigma_b1 = 5,  
  sigma_b2 = 5,  
  mu_alpha = 0,  
  sigma_alpha = 1  
)
```

Calibrating β_1 and β_2 (3)

- ▶ We can now simulate $\tilde{\pi}^s$!

```
prior_sample <- sampling(logi_reg_prior,  
                        data = dat_list,  
                        algorithm = "Fixed_param")
```

- ▶ and extract the $N \times S$ matrix of simulated probability values:

```
prior_sample_fit_pi <- prior_sample %>% spread_draws(prob[condition])
```

- ▶ Notice that this look slightly different from the linear regression case study. We will see why in a second.


Calibrating β_1 and β_2 (4)

- ▶ Let us now visualize the simulated π_i^S ;
- ▶ We can use the the functions in the tidybayes package;

```
prior_sample_fit_pi %>%  
  ggplot(aes(x = condition, y = prob)) +  
  geom_hex(bins = 50) +  
  scale_fill_gradient(low = "white", high = "black") +  
  theme_bw() +  
  labs(x = "individual",  
       y = "Probability") +  
  theme(aspect.ratio = 1,  
        axis.text.x = element_blank(),  
        axis.ticks.x = element_blank(),  
        panel.grid = element_blank())
```

Calibrating β_1 and β_2 (4)

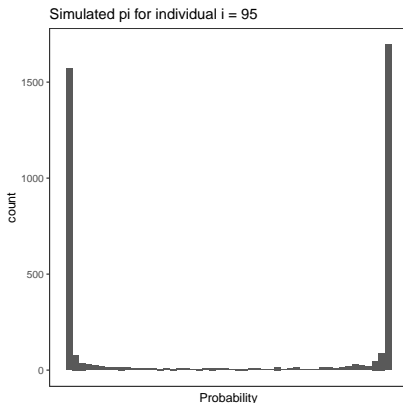
- ▶ Let us now visualize the simulated π_i^S ;
- ▶ We can use the the functions in the `tidybayes` package;



- ▶ This plot is sort of a bird's-eye view of several sequential histograms.

Calibrating β_1 and β_2 (5)


- For example, picking individual, say, $i = 95$ we have:



- Therefore, $\beta_1 = \beta_2 \sim N(0, 5)$ are very informative in that they automatically set probability values to their extremes.

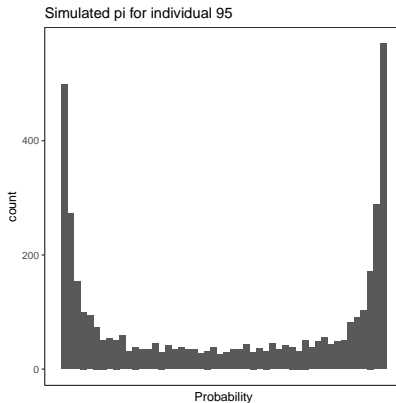
Calibrating β_1 and β_2 (5)

- What if we set $\beta_1 = \beta_2 \sim N(0, 1)$?



Calibrating β_1 and β_2 (6)

- For individual, say, $i = 95$ we have:



- Still very informative!
- So, what can we do if we want to keep prior information limited?

Default priors in logistic regression (1)

- ▶ One again, the problem is the scale-dependence of these coefficients;
- ▶ Even setting a relatively tight prior on β_p , when we multiply these coefficients by x_p , if the range of the corresponding covariate is too big, this will have a disproportionate importance in pushing π to the extremes;
- ▶ One again, we can solve this issue by **standardizing**!
- ▶ In the context of logistic regression this practice is a bit less intuitive, although still fairly straightforward to carry out;
- ▶ The literature suggests that numeric variables should have mean zero and standard deviation one, while categorical variables should be centered on their mean:

$$\dot{z} = \begin{cases} (z - \hat{z})/\text{sd}(z) * 0.5 & \text{if } z \text{ is numeric,} \\ z - \bar{z} & \text{if } z \text{ is binary.} \end{cases}$$

Default priors in logistic regression (2)

- Let us re-define our data list:


```
x1_sc <- as.vector(scale(x1)*0.5)
x2_sc <- as.vector(scale(x2)*0.5)

dat_list <- list(
  N = length(y),
  x1 = x1_sc,
  x2 = x2_sc,
  y = y,
  mu_b1 = 0,
  mu_b2 = 0,
  sigma_b1 = 1,
  sigma_b2 = 1,
  mu_alpha = 0,
  sigma_alpha = 1
)
```

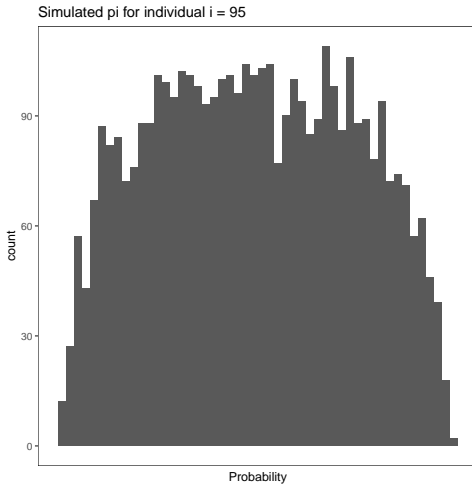
Default priors in logistic regression (3)

```
prior_sample <- sampling(logi_reg_prior,  
                        data = dat_list,  
                        algorithm = "Fixed_param")  
prior_sample_fit_pi <- prior_sample %>% spread_draws(prob[condition])  
  
prior_sample_fit_pi %>%  
  ggplot(aes(x = condition, y = prob)) +  
  geom_hex(bins = 50) +  
  scale_fill_gradient(low = "white", high = "black") +  
  theme_bw() +  
  labs(x = "individual",  
       y = "Probability") +  
  theme(aspect.ratio = 1,  
        axis.text.x = element_blank(),  
        axis.ticks.x = element_blank(),  
        panel.grid = element_blank())
```

Default priors in logistic regression (4)



Default priors in logistic regression (4)



Estimation

- ▶ Once all the priors have been setup, we can use Stan to sample from the implied joint posterior distribution:

$$f(\alpha, \beta_1, \beta_2 | \mathbf{y}, X_1 = x_1, X_2 = x_2)$$

- ▶ To do this, we need to load the second program:

```
logi_reg <- stan_model("logistic_regression.stan")
```

- ▶ Fitting the model to the data now only requires a slightly different sampling statement:

```
fit <- sampling(logi_reg,  
               data = dat_list,  
               chains = 4,  
               cores = 4)  
sample_fit <- rstan::extract(fit)
```

Exploring the results (1)

To quickly visualize the results, we can call the function `summary` as we would do with the standard `lm` output:

```
params <- c("alpha", "beta1", "beta2")  
summary(fit, pars = params)$summary[,1:3]
```

```
##              mean      se_mean      sd  
## alpha  1.1697060 0.004120061 0.2003817  
## beta1  3.1343891 0.009065092 0.4662018  
## beta2 -0.7557881 0.006402164 0.3397158
```

```
params <- c("alpha", "beta1", "beta2")  
summary(fit, pars = params)$summary[,4:8]
```




```
##           2.5%      25%      50%      75%      97.5%  
## alpha  0.7846771 1.0353133 1.166372 1.2978965 1.5599222  
## beta1  2.2340592 2.8213259 3.119159 3.4400835 4.0990099  
## beta2 -1.4233458 -0.9789671 -0.748847 -0.5217524 -0.1039808
```


Exploring the results (2)

- Because of standardization, we need to transform all the coefficients back to the original scale of the data;

```
beta1 <- sample_fit$beta1 * (0.5/sd(x1))  
beta2 <- sample_fit$beta2 * (0.5/sd(x2))  
alpha <- sample_fit$alpha - (beta1*mean(x1)) - (beta2*mean(x2))
```

```
par(mfrow=c(1,3), pty="s")  
hist(beta1)  
hist(beta2)  
hist(alpha)
```



```
par(mfrow=c(1,1), pty="m")
```

Exploring the results (3)

```
summary(beta1)
```

```
##      Min. 1st Qu.  Median      Mean 3rd Qu.      Max.
##  0.3958  0.7478  0.8268  0.8308  0.9119  1.2478
```

```
summary(beta2)
```

```
##      Min. 1st Qu.  Median      Mean 3rd Qu.      Max.
## -0.9867 -0.4968 -0.3800 -0.3835 -0.2648  0.1220
```

```
summary(alpha)
```

```
##      Min. 1st Qu.  Median      Mean 3rd Qu.      Max.
## -0.1763  0.5478  0.7287  0.7248  0.8963  1.7526
```

What if we used glm() instead?

```
summary(glm(y ~ x1 + x2, family = binomial(link = "logit")))


##
## Call:
## glm(formula = y ~ x1 + x2, family = binomial(link = "logit"))
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)   0.7466     0.2886   2.587  0.00967 **
## x1            1.1288     0.1777   6.354  2.1e-10 ***
## x2           -0.4242     0.1968  -2.155  0.03116 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 247.64  on 199  degrees of freedom
## Residual deviance: 157.42  on 197  degrees of freedom
## AIC: 163.42
##
## Number of Fisher Scoring iterations: 6
```

Posterior probabilities (1)

- ▶ Just as in the linear regression case, we can calculate probabilities from the marginal posterior of the estimated parameters:

```
thr <- -0.5  
prob <- mean(beta2>thr); prob
```

```
## [1] 0.75625
```



Credible Intervals

To calculate CrI, we simply take the corresponding quantile values:

```
rbind(quantile(alpha, c(.025, .975)),  
      quantile(beta1, c(.025, .975)),  
      quantile(beta2, c(.025, .975)))
```

```
##           2.5%          97.5%  
## [1,]  0.2142146  1.23708005  
## [2,]  0.5921740  1.08650981  
## [3,] -0.7222782 -0.05276515
```

Model checking (internal consistency - 1)

- ▶ Unlike linear regression, we cannot evaluate goodness of fit by studying min, max and sd of the Posterior Predictive Distribution (PoPD);
- ▶ Instead we need to find a metric that makes sense when using binary data;
- ▶ One such metric is the Correct Classification Rate (CCR), which indicates the proportion of observations correctly classified to either zero or one:

```
ccr <- numeric(length = nrow(sample_fit$y_pred))
for(s in 1:nrow(sample_fit$y_pred)) {


  conf_mat <- table(sample_fit$y_pred[s,], y)
  corr_class <- diag(conf_mat)

  ccr[s] <- (sum(corr_class) / sum(conf_mat)) * 100

}
```

Model checking (internal consistency - 2)

```
par(pty="s")  
hist(ccr, main = "Post. Distribution CCR",  
     xlab="CCR")
```



```
par(pty="m")
```